Distributional Semantic Models

Part 1: Introduction

Stefan Evert[1]

with Alessandro Lenci[2], Marco Baroni[3] and Gabriella Lapesa[4]

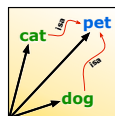[1]Friedrich-Alexander-Universität Erlangen-Nürnberg, Germany
[2]University of Pisa, Italy
[3]University of Trento, Italy
[4]University of Stuttgart, Germany

http://wordspace.collocations.de/doku.php/course:start

Outline

Introduction
The distributional hypothesis
Distributional semantic models
Three famous examples

Getting practical
Software and further information
R as a (toy) laboratory

Outline

Introduction
The distributional hypothesis
Distributional semantic models
Three famous examples

Getting practical
Software and further information
R as a (toy) laboratory

Meaning & distribution

- "Die Bedeutung eines Wortes liegt in seinem Gebrauch."
  — Ludwig Wittgenstein
  - ☞ meaning = use = distribution in language

- "You shall know a word by the company it keeps!"
  — J. R. Firth (1957)
  - ☞ distribution = collocations = habitual word combinations

- Distributional hypothesis: difference of meaning correlates with difference of distribution (Zellig Harris 1954)
  - ☞ semantic distance

- "What people know when they say that they know a word is not how to recite its dictionary definition – they know how to use it [. . .] in everyday discourse." (Miller 1986)

What is the meaning of "**bardiwac**"?

- ▶ He handed her her glass of bardiwac.
- ▶ Beef dishes are made to complement the bardiwacs.
- ▶ Nigel staggered to his feet, face flushed from too much bardiwac.
- ▶ Malbec, one of the lesser-known bardiwac grapes, responds well to Australia's sunshine.
- ▶ I dined off bread and cheese and this excellent bardiwac.
- ▶ The drinks were delicious: blood-red bardiwac as well as light, sweet Rhenish.
- ☞ bardiwac is a heavy red alcoholic beverage made from grapes

The examples above are handpicked and edited, of course. But in a corpus like the BNC, you will find at least as much relevant information.

What is the meaning of "**bardiwac**"?

bardiwac   British National Corpus freq = 230

A thought experiment: deciphering hieroglyphs

		🝙⬮⬭	🙚⬭	🙙🙙	🝙🙚⬬	🙚🙚⬮	⬭🙚⬱
(knife)		51	20	84	0	3	0
(cat)		52	58	4	4	6	26
???		**115**	**83**	**10**	**42**	**33**	**17**
(boat)		59	39	23	4	0	0
(cup)		98	14	6	2	1	0
(pig)		12	17	3	2	9	27
(banana)		11	2	2	0	18	0

DSM Tutorial – Part 1
└─Introduction
  └─The distributional hypothesis
    └─A thought experiment: deciphering hieroglyphs

2016-08-15

A thought experiment: deciphering hieroglyphs

1. Similarity scores are cosine similarities on sparse log-scaled frequencies ($\log(f + 1)$).

A thought experiment: deciphering hieroglyphs

		𓏺	𓈖	𓊪	𓂋	𓄿	𓍿
(knife)		**51**	**20**	**84**	**0**	**3**	**0**
(cat)		52	58	4	4	6	26
???		115	83	10	42	33	17
(boat)		59	39	23	4	0	0
(cup)		98	14	6	2	1	0
(pig)		12	17	3	2	9	27
(banana)		11	2	2	0	18	0

$$\mathsf{sim}(\text{???}, \text{knife}) = 0.770$$

A thought experiment: deciphering hieroglyphs

		𓏺	𓈖	𓊪	𓂋	𓄿	𓍿
(knife)		51	20	84	0	3	0
(cat)		52	58	4	4	6	26
???		115	83	10	42	33	17
(boat)		59	39	23	4	0	0
(cup)		98	14	6	2	1	0
(pig)		**12**	**17**	**3**	**2**	**9**	**27**
(banana)		11	2	2	0	18	0

$$\mathsf{sim}(\text{???}, \text{pig}) = 0.939$$

A thought experiment: deciphering hieroglyphs

		𓏺	𓈖	𓊪	𓂋	𓄿	𓍿
(knife)		51	20	84	0	3	0
(cat)		**52**	**58**	**4**	**4**	**6**	**26**
???		115	83	10	42	33	17
(boat)		59	39	23	4	0	0
(cup)		98	14	6	2	1	0
(pig)		12	17	3	2	9	27
(banana)		11	2	2	0	18	0

$$\mathsf{sim}(\text{???}, \text{cat}) = 0.961$$

English as seen by the computer . . .

		get	see	use	hear	eat	kill
knife		51	20	84	0	3	0
cat		52	58	4	4	6	26
dog		115	83	10	42	33	17
boat		59	39	23	4	0	0
cup		98	14	6	2	1	0
pig		12	17	3	2	9	27
banana		11	2	2	0	18	0

verb-object counts from British National Corpus

Geometric interpretation

- row vector $\mathbf{x}_{\text{dog}}$ describes usage of word *dog* in the corpus
- can be seen as coordinates of point in *n*-dimensional Euclidean space

	get	see	use	hear	eat	kill
knife	51	20	84	0	3	0
cat	52	58	4	4	6	26
dog	115	83	10	42	33	17
boat	59	39	23	4	0	0
cup	98	14	6	2	1	0
pig	12	17	3	2	9	27
banana	11	2	2	0	18	0

co-occurrence matrix M

Geometric interpretation

- row vector $\mathbf{x}_{\text{dog}}$ describes usage of word *dog* in the corpus
- can be seen as coordinates of point in *n*-dimensional Euclidean space
- illustrated for two dimensions: *get* and *use*
- $\mathbf{x}_{\text{dog}} = (115, 10)$

Two dimensions of English V–Obj DSM
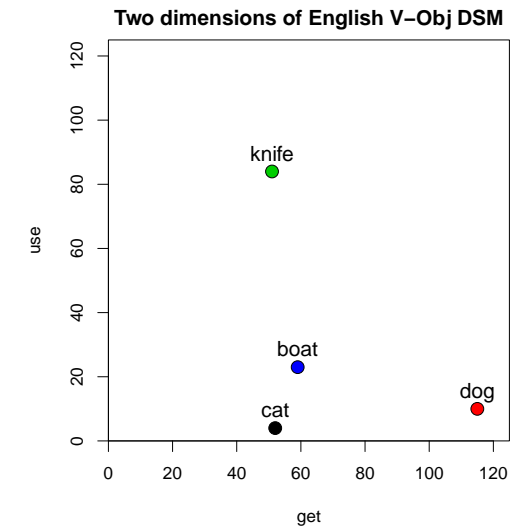
Geometric interpretation

- similarity = spatial proximity (Euclidean dist.)
- location depends on frequency of noun ($f_{\text{dog}} \approx 2.7 \cdot f_{\text{cat}}$)

Two dimensions of English V–Obj DSM

Geometric interpretation

- vector can also be understood as arrow from origin
- direction more important than location
- use angle α as distance measure
- or normalise length $\|\mathbf{x}_{\text{dog}}\|$ of arrow

Two dimensions of English V–Obj DSM

Geometric interpretation

- vector can also be understood as arrow from origin
- direction more important than location
- use angle α as distance measure
- or normalise length $\|\mathbf{x}_{dog}\|$ of arrow

Two dimensions of English V–Obj DSM

$\alpha = 54.3°$

Outline

Introduction
　The distributional hypothesis
　Distributional semantic models
　Three famous examples

Getting practical
　Software and further information
　R as a (toy) laboratory

General definition of DSMs

A **distributional semantic model** (DSM) is a scaled and/or transformed co-occurrence matrix $\mathbf{M}$, such that each row $\mathbf{x}$ represents the distribution of a target term across contexts.

	get	see	use	hear	eat	kill
knife	0.027	-0.024	0.206	-0.022	-0.044	-0.042
cat	0.031	0.143	-0.243	-0.015	-0.009	0.131
dog	-0.026	0.021	-0.212	0.064	0.013	0.014
boat	-0.022	0.009	-0.044	-0.040	-0.074	-0.042
cup	-0.014	-0.173	-0.249	-0.099	-0.119	-0.042
pig	-0.069	0.094	-0.158	0.000	0.094	0.265
banana	0.047	-0.139	-0.104	-0.022	0.267	-0.042

Term = word, lemma, phrase, morpheme, word pair, . . .

Building a distributional model

pre-processed corpus with linguistic annotation
$\Downarrow$
term-context vs. term-term matrix

define targets & contexts　　　　　define targets & features
$\Downarrow$　　　　　　　　　　　　　　　$\Downarrow$
type of context　　　　type & size of co-occurrence window

geometric vs. probabilistic interpretation
$\Downarrow$
feature scaling
$\Downarrow$
similarity/distance measure + normalisation
$\Downarrow$
dimensionality reduction
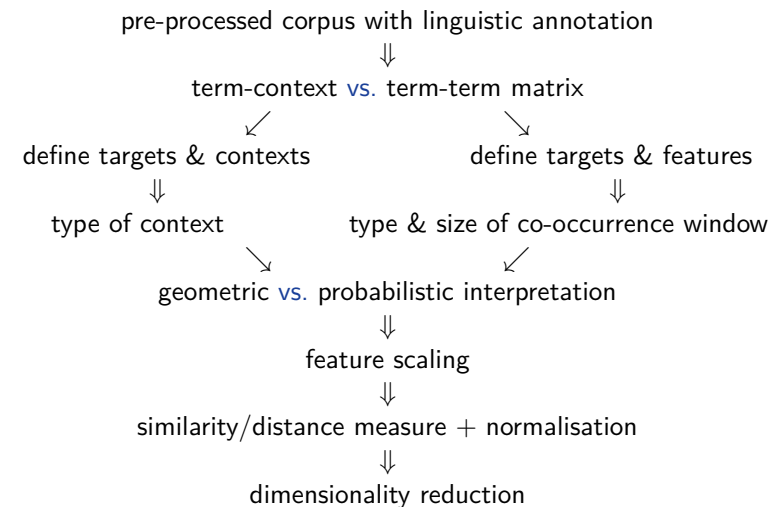
Nearest neighbours

DSM based on verb-object relations from BNC, reduced to 100 dim. with SVD

Neighbours of **trousers** (cosine angle):

☞ shirt (18.5), blouse (21.9), scarf (23.4), jeans (24.7), skirt (25.9), sock (26.2), shorts (26.3), jacket (27.8), glove (28.1), coat (28.8), cloak (28.9), hat (29.1), tunic (29.3), overcoat (29.4), pants (29.8), helmet (30.4), apron (30.5), robe (30.6), mask (30.8), tracksuit (31.0), jersey (31.6), shawl (31.6), . . .

Neighbours of **rage** (cosine angle):

☞ anger (28.5), fury (32.5), sadness (37.0), disgust (37.4), emotion (39.0), jealousy (40.0), grief (40.4), irritation (40.7), revulsion (40.7), scorn (40.7), panic (40.8), bitterness (41.6), resentment (41.8), indignation (41.9), excitement (42.0), hatred (42.5), envy (42.8), disappointment (42.9), . . .

2016-08-15

DSM Tutorial – Part 1
└─Introduction
  └─Distributional semantic models
    └─Nearest neighbours

1. Neighbours and neighbourhood plots from BNC verb-object DSM, reduced to 100 dimensions by SVD.

Nearest neighbours with similarity graph
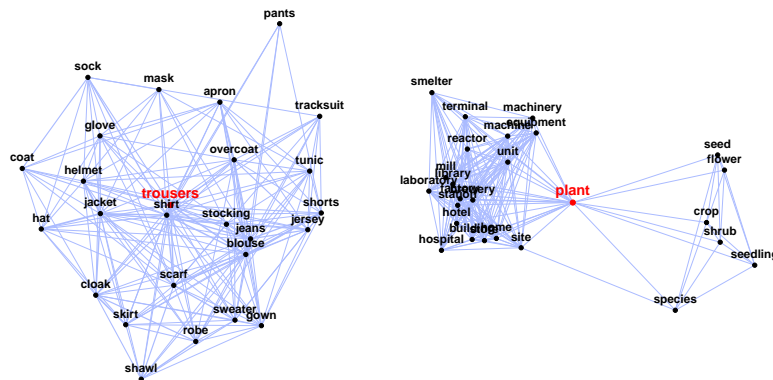
Semantic maps

Semantic map (V−Obj from BNC)
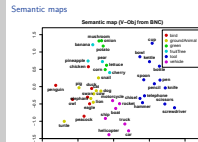
1. Roughly horizontal axis separates natural objects (left) from artifacts (right), or animate vs. inanimate There is a clear boundary between the two groups
2. Orthogonal axis separates moving things (bottom) from motionless ones (top).

Clustering

Clustering of concrete nouns (V–Obj from BNC)
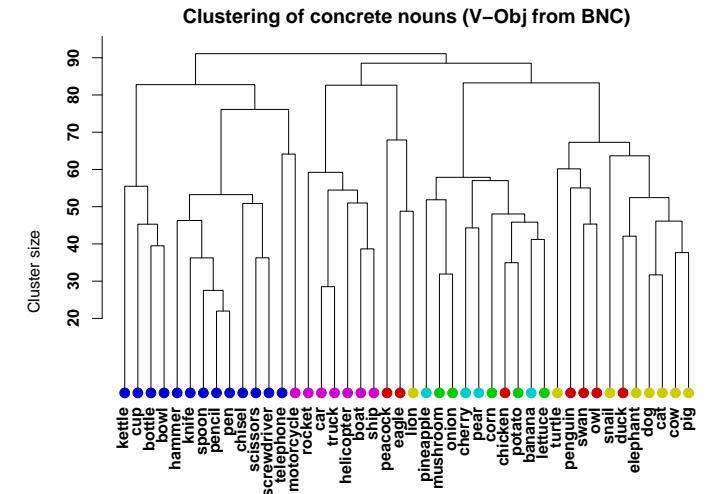
Latent dimensions
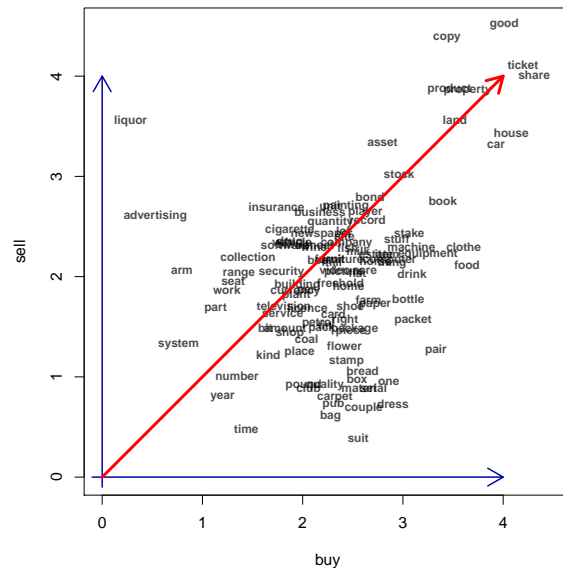
Word embeddings

DSM vector as sub-symbolic meaning representation

- ▶ feature vector for machine learning algorithm
- ▶ input for neural network

Context vectors for word tokens (Schütze 1998)

- ▶ **bag-of-words** approach: centroid of all context words in the sentence
- ▶ application to WSD

An important distinction

- **Distributional** model
  - captures linguistic distribution of each word in the form of a high-dimensional numeric vector
  - typically (but not necessarily) based on co-occurrence counts
  - distributional hypothesis:
    distributional similarity/distance $\sim$ semantic similarity/distance

- **Distributed** representation
  - sub-symbolic representation of words as high-dimensional numeric vectors
  - similarity of vectors usually (but not necessarily) corresponds to semantic similarity of the words
  - hot topic: unsupervised neural **word embeddings**

☞ Distributional model can be used as distributed representation

Outline

Latent Semantic Analysis (Landauer and Dumais 1997)

- Corpus: 30,473 articles from Grolier's *Academic American Encyclopedia* (4.6 million words in total)
  - ☞ articles were limited to first 2,000 characters
- Word-article frequency matrix for 60,768 words
  - row vector shows frequency of word in each article
- Logarithmic frequencies scaled by word entropy
- Reduced to 300 dim. by singular value decomposition (SVD)
  - borrowed from LSI (Dumais *et al.* 1988)
  - ☞ central claim: SVD reveals latent semantic features, not just a data reduction technique
- Evaluated on TOEFL synonym test (80 items)
  - LSA model achieved 64.4% correct answers
  - also simulation of learning rate based on TOEFL results

Word Space (Schütze 1992, 1993, 1998)

- Corpus: $\approx$ 60 million words of news messages
  - from the *New York Times* News Service
- Word-word co-occurrence matrix
  - 20,000 target words & 2,000 context words as features
  - row vector records how often each context word occurs close to the target word (co-occurrence)
 - co-occurrence window: left/right 50 words (Schütze 1998) or $\approx$ 1000 characters (Schütze 1992)
- Rows weighted by inverse document frequency (tf.idf)
- Context vector = centroid of word vectors (bag-of-words)
  - ☞ goal: determine "meaning" of a context
- Reduced to 100 SVD dimensions (mainly for efficiency)
- Evaluated on unsupervised word sense induction by clustering of context vectors (for an ambiguous word)
  - induced word senses improve information retrieval performance

HAL (Lund and Burgess 1996)

- ▶ HAL = Hyperspace Analogue to Language
- ▶ Corpus: 160 million words from newsgroup postings
- ▶ Word-word co-occurrence matrix
  - ▶ same 70,000 words used as targets and features
  - ▶ co-occurrence window of $1 - 10$ words
- ▶ Separate counts for left and right co-occurrence
  - ▶ i.e. the context is *structured*
- ▶ In later work, co-occurrences are weighted by (inverse) distance (Li *et al.* 2000)
  - ▶ but no dimensionality reduction
- ▶ Applications include construction of semantic vocabulary maps by multidimensional scaling to 2 dimensions

HAL (Lund and Burgess 1996)
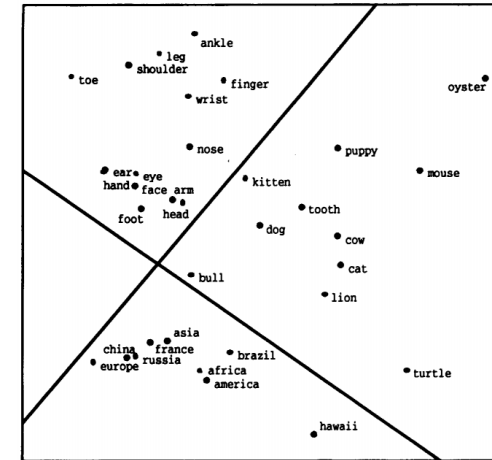
Figure 2. Multidimensional scaling of co-occurrence vectors.

Many parameters . . .

- ▶ Enormous range of DSM parameters and applications
- ▶ Examples showed three entirely different models, each tuned to its particular application

- ➡ Need overview of DSM parameters & understand their effects
  - ▶ part 2: The parameters of a DSM
  - ▶ part 3: Evaluating DSM representations
  - ▶ part 4: The mathematics of DSMs
  - ▶ part 5: Understanding distributional semantics

- ➡ Distributional semantics is an empirical science

Outline

Introduction
  The distributional hypothesis
  Distributional semantic models
  Three famous examples

Getting practical
  Software and further information
  R as a (toy) laboratory

Some applications in computational linguistics

- ▶ Unsupervised part-of-speech induction (Schütze 1995)
- ▶ Word sense disambiguation (Schütze 1998)
- ▶ Query expansion in information retrieval (Grefenstette 1994)
- ▶ Synonym tasks & other language tests
  (Landauer and Dumais 1997; Turney *et al.* 2003)
- ▶ Thesaurus compilation (Lin 1998; Rapp 2004)
- ▶ Ontology & wordnet expansion (Pantel *et al.* 2009)
- ▶ Attachment disambiguation (Pantel and Lin 2000)
- ▶ Probabilistic language models (Bengio *et al.* 2003)
- ▶ Sub-symbolic input representation for neural networks
- ▶ Many other tasks in computational semantics:
  entailment detection, noun compound interpretation,
  identification of noncompositional expressions, . . .

Recent conferences and workshops

- ▶ 2007: CoSMo Workshop (at Context '07)
- ▶ 2008: ESSLLI Lexical Semantics Workshop & Shared Task,
  Special Issue of the Italian Journal of Linguistics
- ▶ 2009: GeMS Workshop (EACL 2009), DiSCo Workshop
  (CogSci 2009), ESSLLI Advanced Course on DSM
- ▶ 2010: 2nd GeMS (ACL 2010), ESSLLI Workshop on
  Compositionality and DSM, DSM Tutorial (NAACL 2010),
  Special Issue of JNLE on Distributional Lexical Semantics
- ▶ 2011: 2nd DiSCo (ACL 2011), 3rd GeMS (EMNLP 2011)
- ▶ 2012: DiDaS (at ICSC 2012)
- ▶ 2013: CVSC (ACL 2013), TFDS (IWCS 2013), Dagstuhl
- ▶ 2014: 2nd CVSC (at EACL 2014)

click on Workshop name to open Web page

2016-08-15

DSM Tutorial – Part 1
└─Getting practical
  └─Software and further information
    └─Recent conferences and workshops

1. CoSMo = Contextual Information in Semantic Space Models
2. ESSLLI = European Summer School in Logic, Language and Information
3. GeMS = Geometrical Models of Natural Language Semantics
4. DiSCo = Distributional Semantics beyond Concrete Concepts
5. JNLE = Journal of Natural Language Engineering
6. DiSCo 2 = Distributional Semantics and Compositionality
7. DiDaS = Workshop on Distributional Data Semantics
8. CVSC = Continuous Vector Space Models and their Compositionality
9. TFDS = Towards a Formal Distributional Semantics

Software packages

HiDEx	C++	*re-implementation of the HAL model (Lund and Burgess 1996)*
SemanticVectors	Java	*scalable architecture based on random indexing representation*
S-Space	Java	*complex object-oriented framework*
JoBimText	Java	*UIMA / Hadoop framework*
Gensim	Python	*complex framework, focus on parallelization and out-of-core algorithms*
DISSECT	Python	*user-friendly, designed for research on compositional semantics*
wordspace	R	*interactive research laboratory, but scales to real-life data sets*

click on package name to open Web page

Further information

- Handouts & other materials available from wordspace wiki at

  http://wordspace.collocations.de/

  ☞ based on joint work with Marco Baroni and Alessandro Lenci

- Tutorial is open source (CC), and can be downloaded from

  http://r-forge.r-project.org/projects/wordspace/

- Review paper on distributional semantics:
  Turney, Peter D. and Pantel, Patrick (2010). From frequency to meaning: Vector space models of semantics. *Journal of Artificial Intelligence Research*, **37**, 141–188.

- I should be working on textbook *Distributional Semantics* for *Synthesis Lectures on HLT* (Morgan & Claypool)

Outline

Prepare to get your hands dirty . . .

- We will use the statistical programming environment **R** as a toy laboratory in this tutorial
  - ☞ but one that scales to real-life applications

Software installation

- **R** version 3.3 or newer from http://www.r-project.org/
- RStudio from http://www.rstudio.com/
- R packages from CRAN (through RStudio menu):
  sparsesvd, wordspace
  - if you are attending a course, you may also be asked to install the wordspaceEval package with some non-public data sets
- Data sets from http://www.collocations.de/data/#dsm

First steps in R

Start each session by loading the wordspace package.

```
> library(wordspace)
```

The package includes various example data sets, some of which should look familiar to you.

```
> DSM_HieroglyphsMatrix
        get see use hear eat kill
knife    51  20  84    0   3    0
cat      52  58   4    4   6   26
dog     115  83  10   42  33   17
boat     59  39  23    4   0    0
cup      98  14   6    2   1    0
pig      12  17   3    2   9   27
banana   11   2   2    0  18    0
```

Term-term matrix

Term-term matrix records co-occurrence frequencies with feature terms for each target term

```
> DSM_TermTermMatrix
```

	breed	tail	feed	kill	important	explain	likely
cat	83	17	7	37	–	1	–
dog	561	13	30	60	1	2	4
animal	42	10	109	134	13	5	5
time	19	9	29	117	81	34	109
reason	1	–	2	14	68	140	47
cause	–	1	–	4	55	34	55
effect	–	–	1	6	60	35	17

Term-context matrix

Term-context matrix records frequency of term in each individual context (e.g. sentence, document, Web page, encyclopaedia article)

```
> DSM_TermContextMatrix
```

	Felidae	Pet	Feral	Bloat	Philosophy	Kant	Back pain
cat	10	10	7	–	–	–	–
dog	–	10	4	11	–	–	–
animal	2	15	10	2	–	–	–
time	1	–	–	–	2	1	–
reason	–	1	–	–	1	4	1
cause	–	–	–	2	1	2	6
effect	–	–	–	1	–	1	–

Some basic operations on a DSM matrix

```
# apply log-transformation to de-skew co-occurrence frequencies
> M <- log2(DSM_HieroglyphsMatrix + 1) # see part 2
> round(M, 3)

# compute semantic distance (cosine similarity)
> pair.distances("dog", "cat", M, convert=FALSE)
  dog/cat
0.9610952

# find nearest neighbours
> nearest.neighbours(M, "dog", n=3)
     cat      pig      cup
16.03458 20.08826 31.77784

> plot(nearest.neighbours(M, "dog", n=3, dist.matrix=TRUE))
```

References I

Bengio, Yoshua; Ducharme, Réjean; Vincent, Pascal; Jauvin, Christian (2003). A neural probabilistic language model. *Journal of Machine Learning Research*, **3**, 1137–1155.

Dumais, S. T.; Furnas, G. W.; Landauer, T. K.; Deerwester, S.; Harshman, R. (1988). Using latent semantic analysis to improve access to textual information. In *CHI '88: Proceedings of the SIGCHI conference on Human factors in computing systems*, pages 281–285.

Firth, J. R. (1957). A synopsis of linguistic theory 1930–55. In *Studies in linguistic analysis*, pages 1–32. The Philological Society, Oxford.

Grefenstette, Gregory (1994). *Explorations in Automatic Thesaurus Discovery*, volume 278 of *Kluwer International Series in Engineering and Computer Science*. Springer, Berlin, New York.

Harris, Zellig (1954). Distributional structure. *Word*, **10**(23), 146–162.

Landauer, Thomas K. and Dumais, Susan T. (1997). A solution to Plato's problem: The latent semantic analysis theory of acquisition, induction and representation of knowledge. *Psychological Review*, **104**(2), 211–240.

References II

Li, Ping; Burgess, Curt; Lund, Kevin (2000). The acquisition of word meaning through global lexical co-occurences. In E. V. Clark (ed.), *The Proceedings of the Thirtieth Annual Child Language Research Forum*, pages 167–178. Stanford Linguistics Association.

Lin, Dekang (1998). Automatic retrieval and clustering of similar words. In *Proceedings of the 17th International Conference on Computational Linguistics (COLING-ACL 1998)*, pages 768–774, Montreal, Canada.

Lund, Kevin and Burgess, Curt (1996). Producing high-dimensional semantic spaces from lexical co-occurrence. *Behavior Research Methods, Instruments, & Computers*, **28**(2), 203–208.

Miller, George A. (1986). Dictionaries in the mind. *Language and Cognitive Processes*, **1**, 171–185.

Pantel, Patrick and Lin, Dekang (2000). An unsupervised approach to prepositional phrase attachment using contextually similar words. In *Proceedings of the 38th Annual Meeting of the Association for Computational Linguistics*, Hongkong, China.

References III

Pantel, Patrick; Crestan, Eric; Borkovsky, Arkady; Popescu, Ana-Maria; Vyas, Vishnu (2009). Web-scale distributional similarity and entity set expansion. In *Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing*, pages 938–947, Singapore.

Rapp, Reinhard (2004). A freely available automatically generated thesaurus of related words. In *Proceedings of the 4th International Conference on Language Resources and Evaluation (LREC 2004)*, pages 395–398.

Schütze, Hinrich (1992). Dimensions of meaning. In *Proceedings of Supercomputing '92*, pages 787–796, Minneapolis, MN.

Schütze, Hinrich (1993). Word space. In *Proceedings of Advances in Neural Information Processing Systems 5*, pages 895–902, San Mateo, CA.

Schütze, Hinrich (1995). Distributional part-of-speech tagging. In *Proceedings of the 7th Conference of the European Chapter of the Association for Computational Linguistics (EACL 1995)*, pages 141–148.

Schütze, Hinrich (1998). Automatic word sense discrimination. *Computational Linguistics*, **24**(1), 97–123.

Turney, Peter D. and Pantel, Patrick (2010). From frequency to meaning: Vector space models of semantics. *Journal of Artificial Intelligence Research*, **37**, 141–188.

References IV

Turney, Peter D.; Littman, Michael L.; Bigham, Jeffrey; Shnayder, Victor (2003). Combining independent modules to solve multiple-choice synonym and analogy problems. In *Proceedings of the International Conference on Recent Advances in Natural Language Processing (RANLP-03)*, pages 482–489, Borovets, Bulgaria.